

Experiment 5

Aim: Perform Regression Analysis using Scipy and Sci-kit learn.

Theory:

1. Linear Regression:

Linear Regression is a supervised learning algorithm used for predicting continuous numerical values. It assumes a linear relationship between the independent variables (features) and the dependent variable (target).

The model minimizes the sum of squared residuals (differences between actual and predicted values) to find the best-fitting line.

It is sensitive to outliers and works best when the data follows a linear trend. Performance is evaluated using metrics like Mean Squared Error (MSE) and R-squared (R^2).

2. Logistic Regression:

Logistic Regression is a classification algorithm used for predicting binary outcomes (0 or 1).

It applies the logistic (sigmoid) function to transform linear outputs into probabilities.

The model predicts a class based on a threshold, typically 0.5.

It uses optimization techniques like gradient descent to minimize the loss function (log loss).

Performance is assessed using accuracy, precision, recall, F1-score, and confusion matrices.

Implementation:

1. Import necessary libraries and load the dataset

```
[ ] import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.metrics import mean_squared_error, r2_score, accuracy_score
```

```
[ ] df = pd.read_csv('diabetes_dataset_with_notes.csv')
```

```
[ ] df.head()
```



	year	gender	age	location	race:AfricanAmerican	race:Asian	race:Caucasian	race:Hispanic	race:Other	hypertension
0	2020	Female	32.0	Alabama	0	0	0	0	1	0
1	2015	Female	29.0	Alabama	0	1	0	0	0	0
2	2015	Male	18.0	Alabama	0	0	0	0	1	0
3	2015	Male	41.0	Alabama	0	0	1	0	0	0
4	2016	Female	52.0	Alabama	1	0	0	0	0	0

```
df.describe()
```

	year	age	race:AfricanAmerican	race:Asian	race:Caucasian
count	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000
mean	2018.360820	41.885856	0.202230	0.200150	0.198760
std	1.345239	22.516840	0.401665	0.400114	0.399069
min	2015.000000	0.080000	0.000000	0.000000	0.000000
25%	2019.000000	24.000000	0.000000	0.000000	0.000000
50%	2019.000000	43.000000	0.000000	0.000000	0.000000
75%	2019.000000	60.000000	0.000000	0.000000	0.000000
max	2022.000000	80.000000	1.000000	1.000000	1.000000

2. Check for missing values and anomalies

```
print("\nMissing values after handling:\n", X.isnull().sum())
```

Missing values after handling:

```
age          0
gender       0
bmi          0
hbA1c_level  0
hypertension 0
heart_disease 0
bmi_age_interaction 0
hbA1c_blood_glucose 0
dtype: int64
```

```
numeric_columns = df.select_dtypes(include=[np.number]).columns
imputer = SimpleImputer(strategy="median")
df[numeric_columns] = imputer.fit_transform(df[numeric_columns])
```

```
df["bmi_age_interaction"] = df["bmi"] * df["age"]
df["hbA1c_blood_glucose"] = df["hbA1c_level"] * df["blood_glucose_level"]
```

```
df = df[df["blood_glucose_level"] < df["blood_glucose_level"].quantile(0.99)]
```

3. Choose your columns and convert categorical values into numerical values for better classification

```
df["gender"] = df["gender"].map({"Male": 0, "Female": 1})
df["smoking_history"] = df["smoking_history"].astype('category').cat.codes
```

```
features = ["age", "gender", "bmi", "hbA1c_level", "hypertension", "heart_disease",
            "bmi_age_interaction", "hbA1c_blood_glucose"]
```

```
X = df[features]
y_reg = df["blood_glucose_level"] # Target for Linear Regression
y_clf = df["diabetes"] # Target for Logistic Regression (Assuming it's binary 0/1)
```

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

4. Train dataset using linear regression

```
[35]
df = df.drop(columns=["location", "clinical_notes", "smoking_history", "gender"])

# Define features (X) and target (y)
X = df.drop(columns=["diabetes"])
y = df["diabetes"]

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize and train linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R-squared Score: {r2}")
```

5. Finding out Mean Squared error and goodness of fit

Mean Squared Error: 0.041312657154867635

R-squared Score: 0.25314264247732976

6. Using Logistic Regression, find out accuracy score

```
X = df.drop(columns=["diabetes"])
y = df["diabetes"]

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize and train logistic regression model
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy}")
print("Classification Report:")
print(report)
```

Accuracy Score:

Accuracy: 0.962550087331758

Classification Report:

	precision	recall	f1-score	support
0.0	0.97	0.99	0.98	18322
1.0	0.82	0.47	0.59	1144
accuracy			0.96	19466
macro avg	0.89	0.73	0.79	19466
weighted avg	0.96	0.96	0.96	19466

Conclusion:

Linear Regression Results:

- The Mean Squared Error (MSE) of 0.0413 indicates that the model's predictions are relatively close to the actual values.

- The R-squared score of 0.2531 shows that the independent variables explain about 25% of the variance in diabetes presence.

Logistic Regression Results:

- The model achieved an accuracy of 96.25%.
- The classification report shows that it performs well for non-diabetic cases (0) with high precision and recall.
- For diabetic cases (1), the recall is 0.47, meaning that while many diabetic cases are correctly identified, some are still being misclassified.